

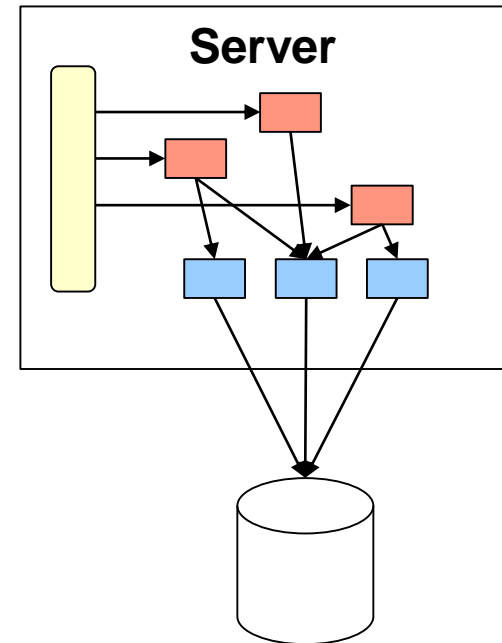
Java Persistence API

- **Persistence Design Patterns**
 - Data Access Object Pattern
 - Service Façade

The screenshot shows a Java Swing application window titled "EAF Lab". It has a menu bar with "File" and "Help". Below the menu bar are four tabs: "Rent Movie", "Movies", "Users", and "Rentals". The "Movies" tab is selected, displaying a table with the following columns: "Movie ID", "Title", "Release Date", "Is Rented?", and "Price Category". The table contains 10 rows of movie data. Below the table is a form with fields for "Movie Title:", "Release Date:", and "Price category:" (with a dropdown menu set to "Children"). At the bottom of the form are buttons for "Refresh", "Cancel", "New ...", "Edit", "Delete", and "Save".

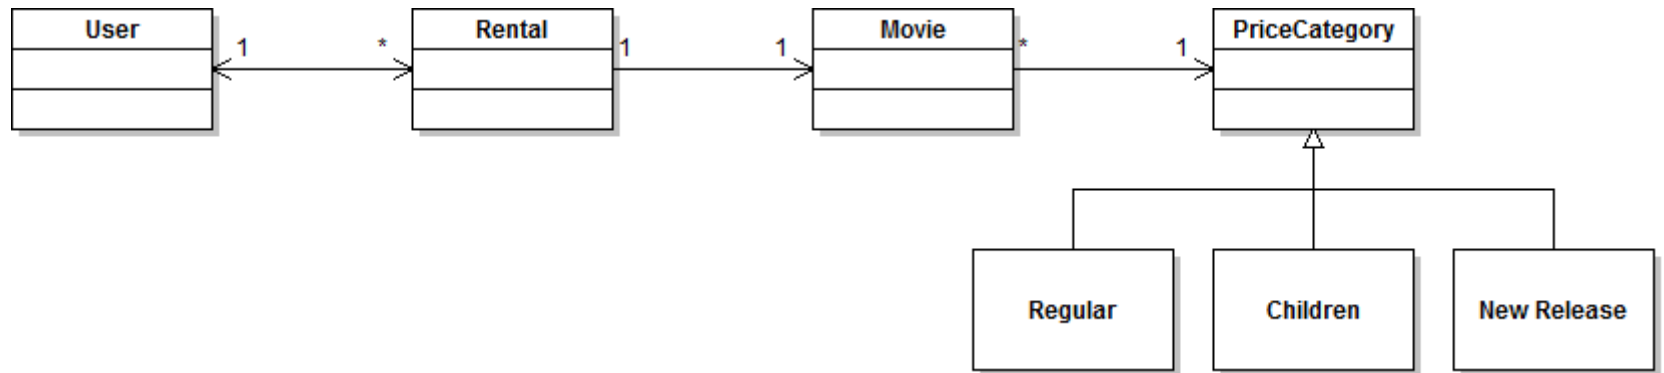
Movie ID	Title	Release Date	Is Rented?	Price Category
1	Marie Curie	2017-06-02	<input checked="" type="checkbox"/>	Regular
2	Curdill	2017-09-20	<input checked="" type="checkbox"/>	Regular
3	The Boss Baby	2017-08-03	<input checked="" type="checkbox"/>	Children
4	Pirates of the Carl...	2017-10-02	<input type="checkbox"/>	New Release
5	Die göttliche Ordn...	2017-09-21	<input type="checkbox"/>	Regular
6	Loving Vincent	2018-05-25	<input type="checkbox"/>	Regular
7	Fast & Furious 7	2018-08-13	<input type="checkbox"/>	New Release
8	Mono	2018-10-01	<input type="checkbox"/>	New Release
9	Swimming with Men	2018-10-03	<input type="checkbox"/>	New Release
10	Jurassic World	2018-10-22	<input type="checkbox"/>	New Release

REST



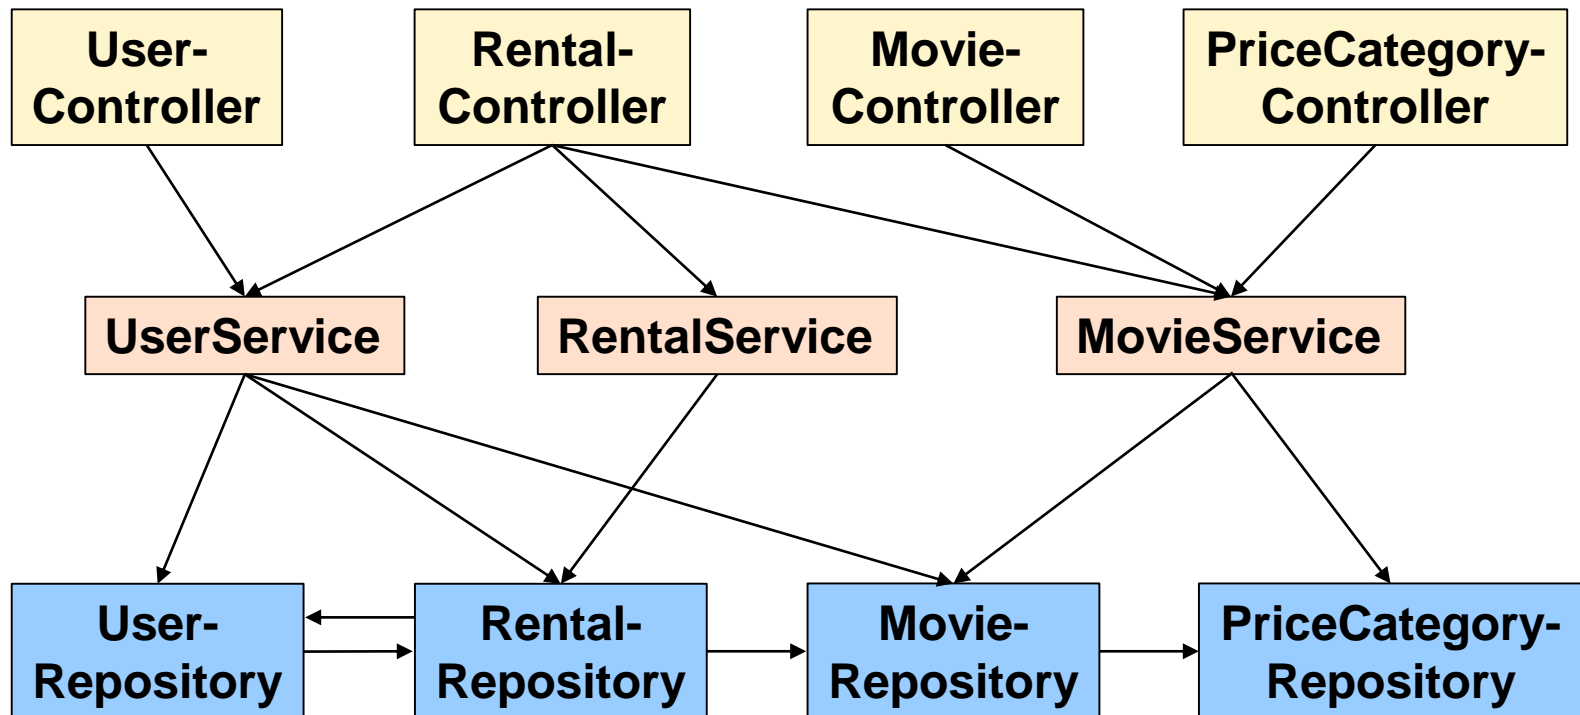
Movie Rental Application

- Model Classes**



Movie Rental Application

- Component Wiring



Service Layer

- **Tasks of the Service Layer**

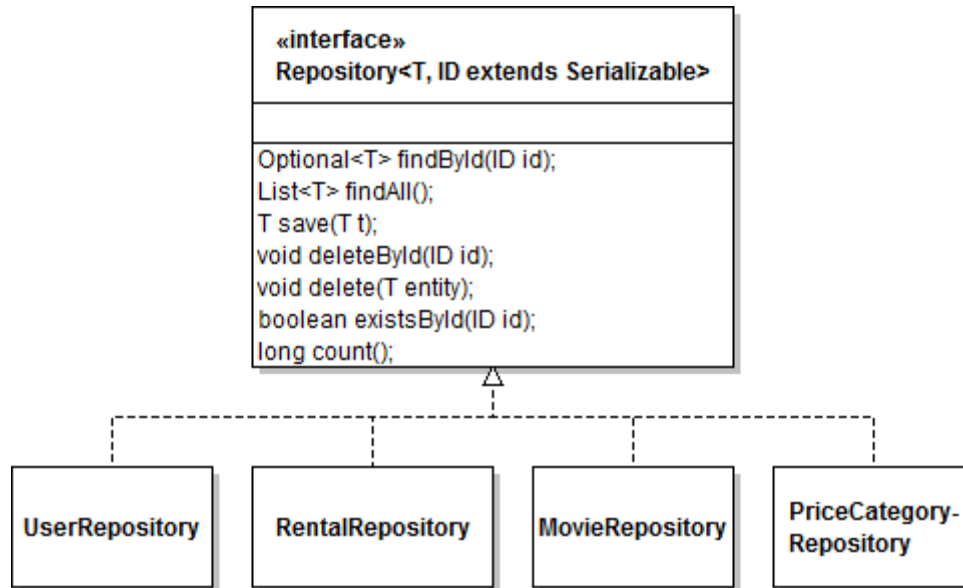
- Core API through which other layers of the application will interface
=> Façade Pattern
- Defines the core business logic, usually calling on one or more DAO methods to achieve this goal
 - Examples:
 - a movie which is returned may be reserved
 - a user which is deleted can have open rentals
 - Some service methods may simply call the corresponding DAO method
- Combines methods defined in the DAOs and assembles them to cohesive business methods that define an atomic unit of work
 - Typically transactional semantics are applied to service methods
 - Spring AOP is used to define transactions

DAO Pattern

- **Data Access Object Pattern (Repository Pattern)**
 - All database access in the system is made through a DAO to achieve encapsulation
 - DAO is also called Repository
 - Each DAO instance is responsible for one primary domain object or entity
 - The DAO is responsible for creations, reads (by primary key), updates, and deletions (CRUD) on the domain object
 - The DAO may allow queries based on criteria other than the primary key (typically *finder methods*) returning collections of the domain object
 - The DAO is *not* responsible for handling transactions, sessions, or connections - these are handled outside the DAO to achieve flexibility

Movie Rental Application

- Repository Interfaces



DAO Pattern

- **Generic DAO / Repository Interface**

```
public interface Repository<T, ID extends Serializable> {  
    Optional<T> findById(ID id);  
    List<T>      findAll();  
  
    T            save(T t);  
  
    void         deleteById(ID id);  
    void         delete(T entity);  
  
    boolean      existsById(ID id);  
    long         count();  
}
```

- Save is used to create and update entries
 - Result is the saved/updated instance

DAO Pattern

- **Separation of Concerns**

- Business & workflow logic
- Persistence issues

Business logic should *not* depend
on chosen persistence technology

- **Strategy Pattern for Data Access**

- Not 100% possible
- DAO type influences its use
 - JDBC-based repositories contains explicit update operations, whereas managed persistence repositories provides object lifecycles, as e.g. JPA, i.e. objects are automatically saved
 - Managed persistence repositories typically allow to define cascade operations on dependent objects, JDBC-based repositories typically not.

DAO Pattern

- **Motivation for the DAO Pattern**

- Encapsulation
- Testability
 - DAO is easier to mock for testing than a hibernate entity manager
- Vendor independence
 - DAO abstracts from different implementations (Hibernate / OpenJPA) of the same DAO Type (JDBC vs. managed)
- Code generation
 - Repository could be created (e.g. using reflection) by frameworks

- **When won't you use the DAO Pattern**

- When the application's business logic consists of data access operations and not much else, it makes sense to include the data access code directly in the business operations.
- JPA Entity Manager or Spring Repositories (Spring Data) **are** DAOs

Java Persistence API

- Persistence Design Patterns
- **JPA Overview**
- JPA Entity Annotations
- JPA Entity Manager
- Spring/SpringBoot & JPA

Java Persistence API

- **Management of Persistence and Object/Relational Mapping**
 - Incorporates contributions from different technologies (communities)
 - Hibernate / TopLink (Oracle) / JDO
- **Standard**
 - JPA 1.0: JSR 220 EJB 3.0 11.05.06 Final Release
 - JPA 2.0: JSR 317 Java Persistence 2.0 10.12.09 Final Release
 - JPA 2.1: JSR 338 Java Persistence 2.1 22.04.13 Final Release
 - JPA 2.2: JSR 338 Java Persistence 2.2 04.08.17 Maintenance Rel.
 - Jakarta Persistence 2.2 Sept 2019 (no changes)
<https://github.com/eclipse-ee4j/jpa-api>
 - Jakarta Persistence 3.0 Sept 2020
<https://jakarta.ee/specifications/persistence/3.0/jakarta-persistence-spec-3.0.pdf>
 - `javax.persistence` => `jakarta.persistence`

JPA Implementations

- **EclipseLink JPA 3.0.2** [JPA 3.0]
 - <http://www.eclipse.org/eclipselink/>
 - Reference Implementation
 - Main Sponsor: Oracle
- **Hibernate 5.5.7** [JPA 3.0]
 - <http://www.hibernate.org/>
 - Main Sponsor: Red Hat
- **Apache OpenJPA 3.1.2** [JPA 2.2]
 - <http://openjpa.apache.org/>
 - Main Sponsor: IBM

We will work in the module eaf with JPA 2.2, i.e. we use classes from the package `javax.persistence`.

JPA Features

- **Features of JPA**

- JPA allows the developer to work directly with objects rather than with SQL statements
- JPA keeps entities in a cache (Persistence Context)
 - Changes on entities are automatically persisted when the transaction is committed
 - No problem with cyclic dependencies
 - No need to implement equals and hashCode
- JPA supports inheritance relationships and supports polymorphic queries
- JPA manages relations
 - Supports automatic loading of dependent objects
 - Supports cascaded delete
 - Supports bidirectional associations

JPA Ingredients

- **Entity Definitions**
 - Metadata configuration
 - Annotations (or XML files)
 - Configuration by Exception
- **Entity Operations**
 - Entity-Manager provides access to the objects (*similar to a Repository*)
 - Query API: find / persist / update / remove
 - JPQL: `SELECT m FROM Movie m WHERE m.title = :title`
 - Persistence context & controlled lifecycle

JPA Entity

@Entity

```
public class Movie {  
    @Id @GeneratedValue  
    private Long id;  
  
    private String title;  
    private boolean rented;  
    private LocalDate releaseDate;  
  
    protected Movie() { }  
    public Movie(String title, LocalDate releaseDate) { ... }  
  
    public String getTitle() { return title; }  
  
    public boolean isRented() { return rented; }  
    public void setRented(boolean rented) { this.rented = rented; }  
    ...  
}
```

MOVIE

ID	TITLE	RENTED	RELEASEDATE

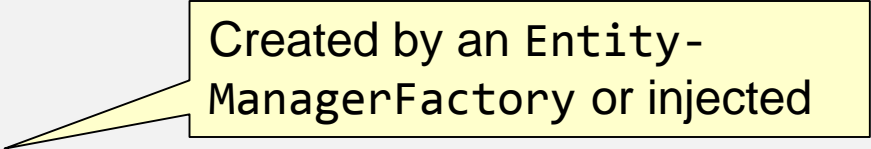
JPA Entity Manager

```
@Service
@Transactional
public class MovieService {

    @PersistenceContext
    private EntityManager em;

    public Long saveNewMovie(String title, LocalDate date) {
        Movie m = new Movie(title, date);
        em.persist(m);
        return m.getId();
    }

    public void rentMovie(Long id) {
        Movie m = em.find(Movie.class, id);
        m.setRented(true); // will be persisted at the end of the TX
    }
}
```



Java Persistence API

- Persistence Design Patterns
- JPA Overview
- **JPA Entity Annotations**
- JPA Entity Manager
- Spring/SpringBoot & JPA

Entity Class: Requirements

- **Entity Class**
 - Entity class must be annotated with the **@Entity** annotation
- **Constraints**
 - Entity class must have a public or protected parameterless constructor
 - Additional constructors may be declared
 - Entity class must not be final
 - Persistent instance fields must be non-final
 - Class must be a top-level class, i.e.
 - not a non-static inner class
 - not an interface
 - not an enum class
 - not a record

Persistent Fields: Supported Types

- **Persistent Fields: Supported Types**
 - Primitive types
 - char, short, int, long, byte, float, double, boolean
 - Serializable Types
 - Strings: *java.Lang.String*
 - Primitive Wrappers: *Integer, Short, Long, Boolean, Double, ...*
 - Big Numericals: *java.math.BigInteger, java.math.BigDecimal*
 - Java temporal types: *java.util.Date, java.util.Calendar*
 - JDBC temporal types: *java.sql.Date, Time, TimeStamp*
 - Java8 temporal types: *java.time.LocalDate, LocalTime, LocalDateTime*
 - Byte- and character arrays: *byte[], char[], Byte[], Character[]*
 - User-defined serializable types (=> DB content is serialized object)
 - Enums
 - **Entity types & Collections of entity types**
 - Collection, Set, List, Map

@Entity

- **Specifies that the class is an entity**

```
@Target(TYPE) @Retention(RUNTIME)
public @interface Entity {
    String name() default "";
}
```

- *name* annotation element
 - This name is used to refer to the entity in queries
 - Must not be a reserved literal in JP-QL (e.g. ORDER)
 - Default: *unqualified name of the entity class*

@Table

- **Specification of primary database table used**

```
@Target(TYPE) @Retention(RUNTIME)
public @interface Table {
    String name() default "";
    String catalog() default "";
    String schema() default "";
    UniqueConstraint[] uniqueConstraints() default {};
}
```

- *name* name of the table *unqualified name of the entity class*
- *catalog* catalog of the table *standard catalog*
- *schema* schema of the table *standard schema*
- *uniqueConstraints* constraints applied to generated DDL tables

```
@Table(name = "DEMO_USER", uniqueConstraints = {
    @UniqueConstraint(columnNames = { "PHONENR", "AREA-CODE" })
})
```

@Column

- **Specification of mapped column for a persistent property**

```
@Target({METHOD, FIELD}) @Retention(RUNTIME)
public @interface Column {
    String name() default "";           // name of column
    boolean unique() default false;    // if DB column is unique
    boolean nullable() default true;   // if DB column is nullable
    boolean insertable() default true; // if manipulation with a
                                        // sql insert is allowed

    boolean updatable() default true;
    String columnDefinition() default ""; // e.g. CLOB / BLOB
    String table() default "";           // table in which field
                                        // is stored (sec. table)

    int length() default 255;          // size for strings
    int precision() default 0;         // decimal precision
    int scale() default 0;             // decimal scale
}
```

@Enumerated

- **Specification of mapping of enumerations**

```
@Target({METHOD, FIELD}) @Retention(RUNTIME)
public @interface Enumerated {
    EnumType value() default ORDINAL;
}
```

- Value field
 - EnumType.ORDINAL
 - Value is stored as an integer
 - EnumType.STRING
 - Value is stored as a string

@Id

- **Specification of primary key field**

```
@Target({METHOD, FIELD}) @Retention(RUNTIME)
public @interface Id {
}
```

- Possible types for primary keys

- Primitive types: *byte, int, short, long, char*
 - Wrapper types: *Byte, Integer, Short, Long, Character*
 - Arrays of primitives or wrappers
 - Strings: *java.lang.String*
 - Large numerics: *java.math.BigInteger*
 - Temporal types: *java.util.Date, java.sql.Date, java.time.LocalDate*
 - UUID: *java.util.UUID* (supported by Hibernate)

Primary Keys: Generation

- **Assigned**
 - Primary keys may be assigned by the application, i.e. no key generation is necessary
 - E.g. language table: primary key is the ISO country code
 - UUID (global unique identifier)
- **Identity**
 - Auto increment supported by some DBs
- **Sequence**
 - Some DBs support sequences which generate unique values (e.g. Oracle, PostgreSQL)
- **Table**
 - Primary keys are stored in a separate PK table

@GeneratedValue Annotation

- **Specification of primary key generation method**

```
@Target({METHOD, FIELD}) @Retention(RUNTIME)
public @interface GeneratedValue {
    GenerationType strategy() default AUTO;
    String generator() default "";
}

public enum GenerationType {TABLE, SEQUENCE, IDENTITY, AUTO};
```

- *strategy* primary key generation strategy
- *generator* name of the generator as specified in SequenceGenerator or TableGenerator annotations

Annotations Example

```
@Entity
@Table(name="EMP")
public class Employee {
    public enum Type {FULL, PART_TIME};

    protected Employee() { }
    public Employee(String name, Type type) {
        this.name = name; this.type = type;
    }

    @Id
    @GeneratedValue(strategy = GenerationType.IDENTITY)
    Long id;

    @Enumerated(EnumType.STRING)
    @Column(name = "EMP_TYPE", nullable = false)
    Type type;

    String name;
}
```

Annotations Example

```
CREATE SEQUENCE PUBLIC.SYSTEM_SEQUENCE_6AD8FFC7
  START WITH 1
  INCREMENT BY 1
  MINVALUE 1;

CREATE MEMORY TABLE PUBLIC.EMP(
  ID BIGINT NOT NULL
    DEFAULT NEXT VALUE FOR PUBLIC.SYSTEM_SEQUENCE_6AD8FFC7
    NULL_TO_DEFAULT SEQUENCE PUBLIC.SYSTEM_SEQUENCE_6AD8FFC7,
  NAME VARCHAR(255),
  EMP_TYPE VARCHAR(255) NOT NULL,
  PRIMARY KEY (ID)
)
```

- IDENTITY is modelled with a sequence which is updated on each insert

Associations

- **Relationship modeling annotations**

- @OneToOne
 - @OneToMany
 - @ManyToOne
 - @ManyToMany
-
- Annotations must be applied to the corresponding property

```
@OneToOne  
Address address;  
@OneToMany  
List<Rental> rentals;
```

- Associations may be **unidirectional** or **bidirectional**
 - For bidirectional associations, one side is marked with "mappedBy" (*inverse side*), the other side is called the *owning side*

ManyToOne: User - Rental: bidirectional

```
@Entity
public class Rental {
    @Id @GeneratedValue
    private int id;

    @ManyToOne          // Rental is the owner of the relationship
    @JoinColumn(name = "USER_FK") // optional
    private User user;

    public Rental() { }

    public User getUser() { return user; }
    public void setUser (User user) { this.user = user; }

    public int getId() { return id; }
    public void setId(int id) { this.id = id; }
}
```

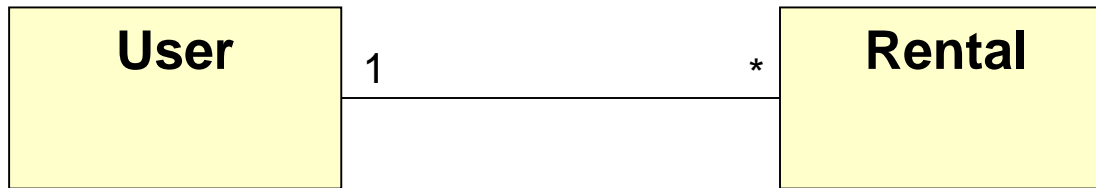
ManyToOne: User - Rental: bidirectional

```
@Entity
public class User {
    ...
    @OneToMany(mappedBy = "user")
    private List<Rental> rentals;
        // this is the inverse side of the relationship

    public Collection<Rental> getRentals() {
        return rentals;
    }
    public void setRentals(Collection<Rental> rentals) {
        this.rentals = rentals;
    }
}
```

ManyToOne: User - Rental: bidirectional

- Representation in Database



```
create table user (  
  ID integer,  
  FIRSTNAME varchar(32),  
  LASTNAME varchar(32),  
);
```

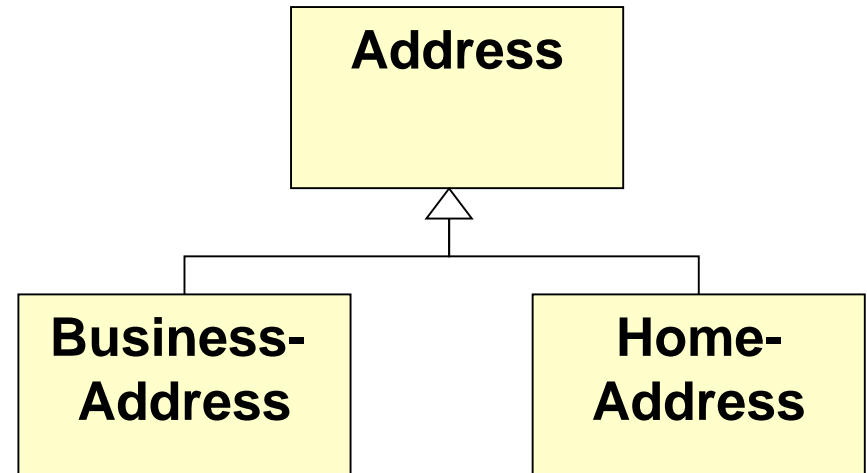
```
create table rental (  
  ID integer,  
  USER_FK integer,  
  RENTED_DAYS integer,  
  RENTED_DATE date  
);
```

An arrow points from the 'USER_FK' field in the 'rental' table definition to the 'FIRSTNAME' field in the 'user' table definition, indicating a foreign key relationship.

Inheritance

- **Entities**

- All classes have to be declared to be an @Entity classes



- **Representation**

- All classes are (by default) stored in a SINGLE TABLE

- **Annotation**

- @DiscriminatorColumn(name="PRICECATEGORY_TYPE")
 - defines the name of the column where dynamic type is stored
 - to be defined on the base class
- @DiscriminatorValue("Children")
 - to be defined on the concrete base- and subclasses

Java Persistence API

- Persistence Design Patterns
- JPA Overview
- JAP Entity Annotations
- **JPA Entity Manager**
- Spring/SpringBoot & JPA

Entity Manager

- **EntityManager = API to access the database**
 - Entity manager stores the entities in a persistence context
 - Entity manager manages the lifecycle of entity instances
 - An entity contained in a persistence context is managed (attached)
 - An entity not associated with a persistence context is unmanaged (detached)
- **Persistence Context = set of managed objects**
 - Only one entity instance with the same persistent identity may exist in a Persistence Context
 - Entity manager automatically synchronizes the content of the persistence context with the database
 - Changes made on managed objects are automatically flushed
- **Declaration**

```
@PersistenceContext  
private EntityManager em;
```

Entity Manager

```
public interface EntityManager {  
    public void persist(Object entity);  
    public <T> T merge(T entity);  
    public void remove(Object entity);  
    public <T> T find(Class<T> entityClass, Object primaryKey);  
  
    public void flush();  
    public void setFlushMode(FlushModeType flushMode);  
    public FlushModeType getFlushMode();  
  
    public void refresh(Object entity);  
    public void clear();  
    public boolean contains(Object entity);  
    public void detach(Object entity);  
  
    public Query createQuery(String qlString);  
    public Query createNamedQuery(String name);  
    ...  
}
```

Entity Manager

- **persist**
 - Makes an instance managed and persistent
- **remove**
 - Removes the entity instance from the database
- **find**
 - Finds an entity by its primary key
- **merge**
 - Merges the state of the given entity into the current persistence context
 - Returns new (unique) instance
 - May be used for insert operations as well (managed instance is returned)
- **refresh**
 - Refresh the state of the (managed) instance from the database, discarding changes made on the entity, if any

Entity Manager

- **flush**
 - Synchronize the persistence context to the underlying database
- **contains**
 - Check if the instance belongs to the current persistence context (but not whether the entity is stored on the database)
- **clear**
 - Clear the persistence context, causing all managed entities to become detached
 - Changes made to entities that have not been flushed to the database will not be persisted
- **detach**
 - Removes the given entity from the persistence context

Queries

- **createQuery**

- Creates an instance of Query for executing an JPQL statement
- Returns a Query / TypedQuery<T> object
 - Parameters:
 - `q.setParameter(String, Object)` sets named parameters
 - `q.setParameter(int, Object)` sets positional parameters
 - Results:
 - `q.getResultList()`
 - `q.getSingleResult()`

```
TypedQuery<Movie> q = em.createQuery(  
    "SELECT m FROM Movie m WHERE m.title = :title",  
    Movie.class);  
q.setParameter("title", title);  
List<Movie> movies = q.getResultList();
```

Transactions

- **Transactions**

- A transaction is needed for all EntityManager operations that change anything in the DB
 - No transaction is required for SELECT queries (without locking)

```
javax.persistence.TransactionRequiredException:  
No EntityManager with actual transaction available for current  
thread - cannot reliably process 'merge' call
```

- **Declare your Services as @Transactional**

- Probably tests have to be declared as Transactional as well otherwise entities are detached

Transactions and Transaction strategies will be covered later

Java Persistence API

- Persistence Design Patterns
- JPA Overview
- JPA Entity Annotations
- JPA Entity Manager
- **Spring/SpringBoot & JPA**

Spring Boot

- Dependencies**

```
dependencies {
    implementation("org.springframework.boot:spring-boot-starter-data-jpa")
    runtime("com.h2database:h2")
}
```

Category/License	Group / Artifact	Version	Updates
EDL 1.0 EPL 2.0	 jakarta.persistence » jakarta.persistence-api	2.2.3	3.0.0
EPL 2.0	 jakarta.transaction » jakarta.transaction-api	1.3.3	2.0.0
O/R Mapping LGPL 2.1	 org.hibernate » hibernate-core	5.4.32.Final	6.0.0.Alpha9
AOP Apache 2.0	 org.springframework » spring-aspects	5.3.10	✓
AOP Apache 2.0	 org.springframework.boot » spring-boot-starter-aop	2.5.5	✓
Apache 2.0	 org.springframework.boot » spring-boot-starter-jdbc	2.5.5	✓
Apache 2.0	 org.springframework.data » spring-data-jpa	2.5.5	✓

Spring Boot Properties

- **spring.jpa**

```
org.springframework.boot.autoconfigure.orm.jpa.JpaProperties
```

```
└─ JpaProperties
   └─ JpaProperties()
      ├── properties : Map<String, String>
      ├── mappingResources : List<String>
      ├── databasePlatform : String
      ├── database : Database
      ├── generateDdl : boolean
      ├── showSql : boolean
      └── openInView : Boolean
```

- **spring.jpa.hibernate**

```
org.springframework.boot.autoconfigure.orm.jpa.HibernateProperties
```

```
└─ HibernateProperties
   └─ HibernateProperties()
      ├── naming : Naming
      ├── ddlAuto : String
      └── useNewIdGeneratorMappings : Boolean
```

Spring Boot Properties

- **spring.datasource**

```
org.springframework.boot.autoconfigure.jdbc.DataSourceProperties
```

```
DataSourceProperties
  DataSourceProperties()
  classLoader : ClassLoader
  generateUniqueName : boolean
  name : String
  type : Class<? extends DataSource>
  driverClassName : String
  url : String
  username : String
  password : String
  jndiName : String
  schema : List<String>
  embeddedDatabaseConnection : EmbeddedDatabaseConnection
  xa : Xa
  uniqueName : String
```

- **spring.h2.console**



```
org.springframework.boot.autoconfigure.h2.H2ConsoleProperties
```

```
H2ConsoleProperties
  H2ConsoleProperties()
  path : String
  enabled : boolean
  settings : Settings
```

Spring Boot Properties

- **spring.sql.init**

```
org.springframework.boot.autoconfigure.sql.init.  
    SqlInitializationProperties
```

```
▲  SqlInitializationProperties  
  ●  SqlInitializationProperties()  
    □ schemaLocations : List<String>  
    □ dataLocations : List<String>  
    □ platform : String  
    □ username : String  
    □ password : String  
    □ continueOnError : boolean  
    □ separator : String  
    □ encoding : Charset  
    □ mode : DatabaseInitializationMode
```

IDEs provide auto completion in
application.properties

- **Documentation**

- <https://docs.spring.io/spring-boot/docs/current/reference/html/application-properties.html>

Spring Boot Entity Scanning

- **EntityScan**

- Spring Boot uses "Entity Scanning" to look for Entities
 - The package containing the class annotated with `@EnableAutoConfiguration` is used as base (sub packages are scanned as well)
 - May be overwritten with `@EntityScan`

```
@Target(ElementType.TYPE) @Retention(RetentionPolicy.RUNTIME)
public @interface EntityScan {

    // list of base packages to scan for entities
    String[] basePackages() default {};
    String[] value() default {}; // alias for basePackages

    // type safe variant to specify base packages to scan for
    // entities
    Class<?>[] basePackageClasses() default {};
}
```

Spring Boot Data-Source Properties

- **Datasource Properties**

```
# JDBC URL of the database
spring.datasource.url=jdbc:h2:mem:testdb
spring.datasource.url=jdbc:h2:file:./movierental-db
spring.datasource.url=jdbc:h2:./movierental-db
spring.datasource.url=jdbc:h2://localhost:8888/test
```

- If no URL is specified, Spring Boot auto generates the database name for in-memory databases (H2, HSQL and Derby), name is shown in the log, e.g. `jdbc:h2:mem:94b35e1f-228d-4ab1-99da-bd2879b3605d`

- This behavior can be disabled using

```
spring.datasource.generate-unique-name=false
```

- Data source name is then `testdb` for embedded DBs, null otherwise, can be changed with

```
spring.datasource.name=movie-db
```

Spring Boot Data-Source Properties

- **Login username and password of the database**

```
# username & password (sa/<empty> for H2 in-memory DB)
spring.datasource.username=movieadmin
spring.datasource.password=changeme
```

- **Name of the JDBC driver**

```
# Autodetected based on the URL by default
spring.datasource.driver-class-name=org.h2.Driver
```

- **Using a different connection pool**

```
spring.datasource.type=org.apache.tomcat.jdbc.pool.DataSource
```

- HakariCP (default)
- Tomcat Connection Pooling
- Commons DBCP2
- Oracle UCP

Spring Boot Schema Creation

- **Creating Schemas**

- By default, JPA databases will be created automatically only if an embedded database (H2, HSQL or Derby) is used
- Default behavior can be overridden

- Hibernate property

`hibernate.hbm2ddl.auto`

`spring.jpa.hibernate.ddl-auto=create-drop`

- none
 - validate validate the schema, makes no changes to the database
 - update update the schema
 - create creates the schema, destroying previous data
 - create-drop create and then drop the schema at the end of the session
 - Defaults
 - For an embedded DB the default is create-drop
 - For a non-embedded DB the default is none

Database Initialization using Hibernate

- **import.sql**

- A file named `import.sql` in the root of the classpath is executed on startup if `d11-auto=create` or `create-drop`.
- This is a Hibernate-feature (and has nothing to do with Spring)
- Name of the file can be configured:

```
# configuration for Hibernate (import.sql by default)
javax.persistence.hibernate.hbm2ddl.import_files=d1.sql,d2.sql
```

Database Initialization using Spring JDBC

- **Database Initialization using SQL-Scripts (schema.sql / data.sql)**
 - Configuration properties
 - `spring.sql.init.mode` (always/embedded/never)
 - Defines whether scripts are executed
 - `spring.sql.init.schema-locations/spring.sql.init.data-locations`
 - path to schema and data files; can be absolute, relative or classpath based
 - `spring.sql.init.continue-on-error`
 - Indicates whether program should stop if an error occurs while initializing the database
 - **Default Parameters**

```
spring.sql.init.mode=embedded  
spring.sql.init.schema-locations=classpath:/schema.sql  
spring.sql.init.data-location=classpath:/data.sql  
spring.sql.init.continue-on-error=false
```

Database Initialization using Spring JDBC

- **Initialization Order**

- Script-based DataSource initialization is performed, by default, *before* any JPA EntityManagerFactory beans are created, i.e. before schema is created by JPA
 - Typically either script-based or JPA-based initialization is used
 - JPA overrides schemas defined by scripts
- If scripts should be executed *after* the schema was created by Hibernate, then set the following property:

```
spring.jpa.defer-datasource-initialization=true
```

- `schema.sql` can then be used to make additions to any created schema
- `data.sql` can be used to populate the created schemas

Spring Boot H2 Properties

- H2's Browser Based Web Console**

```
spring.h2.console.enabled=true    # default: false
spring.h2.console.path=/h2       # default: /h2-console, must start
                                # with a "/"
```

English ▼ Preferences Tools Help

Login

Saved Settings: Generic H2 (Embedded) ▼

Setting Name: Generic H2 (Embedded) Save Remove

Driver Class: org.h2.Driver

JDBC URL: jdbc:h2:mem:testdb

User Name: sa

Password:

Connect Test Connection

Spring Boot Additional Properties

- **JPA Logging**

```
# Enable logging of SQL statements, default is false  
spring.jpa.show-sql=true
```

- **Hibernate Properties**

- Additional native properties to be set on a JPA provider (as e.g. Hibernate) start with `spring.jpa.properties.*` (this prefix is dropped before passing them)
- https://docs.jboss.org/hibernate/orm/current/userguide/html_single/Hibernate_User_Guide.html#configurations

```
# JDBC connection transaction isolation level  
spring.jpa.properties.hibernate.connection.isolation=8  
# Allows Hibernate to generate SQL optimized for a particular DB  
# Value is chosen based on the metadata returned by JDBC-driver  
spring.jpa.properties.hibernate.dialect=org.hibernate.dialect.H2Dialect
```